

KEVIN CHEN

Austin, Texas | (915) 330-6646 | kevinchen2662@gmail.com | [linkedin.com/in/kevinchen76](https://www.linkedin.com/in/kevinchen76) | [GitHub](#) | [Portfolio](#)

EDUCATION

The University of Texas at Austin

Austin, TX

B.S. in Economics, Minor in Computer Science, Data Science, Applied Statistics

Graduated May 2025

Coursework: data analytics, databases, data visualizations, data science, software design, probability, econometrics

SKILLS

Programming Languages: Python, SQL, Java, C++, R/R studio, HTML/CSS, STATA

Framework/Tools: AWS, Azure, Snowflake, Databricks, Tableau, Power BI, pandas, scikit-learn, PyTorch, matplotlib, MySQL, Git/GitHub, Excel (Advanced)

Certifications: CFA Level 1 Candidate (February 2027), AWS Data Engineer Associate (In Progress), Microsoft Azure Administrator Associate (In Progress), Databricks Data Engineer Associate, Salesforce Tableau Data Analyst

PROFESSIONAL EXPERIENCE

Urban Information Lab – The University of Texas at Austin

Austin, TX

Research Professional

February 2024 – Present

- Co-authored 13+ peer-reviewed publications on LLM safety in AI ethics, accumulating 77+ citations within 12 months
- Spearheaded Python text-mining pipelines across 4+ code repositories, extracting and classifying hundreds of institutional AI policy documents; defined acceptance criteria and ran regression-style validation checks via Git/GitHub
- Developed Python/SQL workflows for ingestion, cleaning, and schema validation with automated quality checks; maintained data dictionaries, release notes, and testing procedures, cutting manual QA errors ~40%
- Distilled complex model evaluation results into 5+ Tableau dashboards and 50+ Python visualizations, enabling 6-person team to track research milestones and communicate findings to technical and non-technical stakeholders

Higginbotham Insurance

Odessa, TX

Data Analytics Intern

June 2024 – Aug 2024

- Ensured data integrity of 20+ client-facing Tableau dashboards and Excel reports by reconciling outputs against Snowflake data, verifying accuracy and supporting client retention and renewal decisions across 3 account teams
- Optimized 12+ Snowflake SQL queries underpinning client proposal and renewal reporting tables, cutting average execution time by up to 35% and accelerating turnaround on revenue-critical data requests
- Collaborated with account managers to resolve urgent data requests, ensuring on-time delivery of all materials supporting active client engagements throughout a 10-week internship

Texas Public Policy Foundation

Austin, TX

Policy Research Intern

June 2023 – August 2023

- Researched and authored data-backed policy briefs on data privacy, online safety, deepfakes, and AI regulation; synthesized expert roundtables and incorporated stakeholder feedback into tracked revisions under tight deadlines

PROJECTS

Safe-Child-LLM Benchmark | Red-Team Evaluation Suite [\[GitHub\]](#) [\[Paper\]](#)

Austin, TX

Lead Developer

Jun 2025 – Present

- Spearheaded development of a published AI safety evaluation framework benchmarking 5+ AI systems on child-safety risks, informing responsible deployment guidelines for child-facing technology
- Implemented automated testing to validate evaluation consistency across model versions, enabling iterative benchmark expansion across developmental age groups for external research stakeholders, replacing manual scoring

LLM Ethics Benchmark | Moral Reasoning Evaluation Framework [\[GitHub\]](#) [\[Paper\]](#)

Austin, TX

Contributor and Developer

May 2025 – Present

- Contributed to a published moral evaluation framework for AI systems, developing scoring components that standardized multi-dimensional ethical assessment across 4+ LLM models

PUBLICATIONS

- Jiao J., Afroogh S., Chen K., et al. (2026). *Generative AI and LLMs in industry: a text-mining analysis of guidelines across 14 industrial sectors*. Nature Humanities and Social Sciences Communications, 13, 410 [doi](#)
- Jiao J., Afroogh S., Murali A., Chen K., et al. (2025). *LLM Ethics Benchmark: A three-dimensional assessment of moral reasoning in LLMs*. Nature Scientific Reports, 15(1), 34642 [doi](#)
- Jiao J., Afroogh S., Chen K., et al. (2025). *The global landscape of academic guidelines for generative AI and LLMs*. Nature Human Behaviour, 9(4), 638–642 [doi](#)